Assignment 2: The Structured Query Language

The following two sections contain formatting requirements for your homework submission.
Please read both sections; failure to follow these guidelines will result in point deductions.

General Formatting Guidelines

All the rules about good coding style apply to SQL statements as well, because they also need to be maintained over time. Here are some of the most important ones:

- Keep lines to 80 characters or less. (We're serious about this one.)
- Always give aggregates and other computed results a clear and meaningful name when they are referred to by other parts of the query, or when they appear in a view.
- Don't wrap long SQL statements just anywhere! It's best to wrap SQL statements at the start of a major clause, or at the end of a component of a given clause. Make it easy to understand.
- Similarly, follow a uniform indentation style that clearly indicates the various parts of SQL queries. For example, if you have a large number of tables in a FROM clause, or you have a derived relation, make sure everything in the FROM clause is indented to show its extent.
- Make sure indentation is done with *spaces* and not *tabs*, as tabs are not necessarily the same width from machine to machine, or editor to editor.
- You are strongly encouraged to capitalize SQL keywords and use lower-case table/column names, but this is not required. Whatever you do, be <u>consistent</u>.

We have provided a style-checking tool for you to run against your files. While you are not required to run it prior to submitting your SQL, it is highly recommended, as we will be using the same style-checker to deduct points for style errors. You can run it as follows:

```
# Presuming you have done "chmod u+x check.py" first:
check.py file1.sql file2.sql ...

# If you don't want to make check.py executable:
python check.py file1.sql file2.sql ...
```

Submission Formatting Requirements

For this assignment we are going to use a new tool for automating parts of the grading. Your answers for the problems will be put into various SQL files; the specific filenames will be indicated in each section. (You will be provided a template to fill in, so that your life is easier.)

You must submit your entire homework solution as a single archive file named **username-cs121hw2.tgz**, where "username" is your CMS (preferred) or IMSS account username. This will allow us to test your SQL statements easily, without having to copy/paste them out of a PDF document or some other strange format. Please also make sure not to use some strange text file format like Unicode or RTF. (The style checker will catch this if you make such a mistake.)

You also need to follow specific guidelines for naming and formatting your SQL files; the format to use is as follows:

The answer for each problem must begin with a comment like this:

```
-- [Problem 1a]
```

In this case, "1a" is the number of the specific problem you are solving. Each question will specify what to put as the problem number, e.g. "5" or "6e".

Any descriptive comments should follow this annotation, for example:

```
-- [Problem 1a]
-- Comments for this problem.
```

The SQL statement can be anywhere after the problem annotation and the comment, like this:

```
-- [Problem 1a]
-- Comments for this problem.
SELECT * FROM t WHERE a > 4;
```

Of course, you can put blank lines between any of these portions to improve readability. **However, do not put any code before the first** "`-- [Problem ..]`" **annotation! Any code before this annotation will not be run and may cause the auto-grader to mark all your responses as incorrect!**

As stated earlier, an archive containing template files is provided with this assignment, so you can simply fill in the answers between the required annotations. When you complete the assignment, repackage this archive and submit it on the course website.

MySQL Database Access
You can use a database to test virtually all of your answers to this assignment. You should already have a database to use; the connection information will have been sent to you in an email. If you haven't received this connection information, please contact Donnie so he can set one up for you.

You are encouraged to consult the MySQL 5.5 online documentation[1] if you have any questions about syntax or functionality. You will be exploring many of MySQL's features, so it is a good idea to become familiar with the manual now. **We are currently using MariaDB 5.5.33a**, which roughly corresponds to MySQL 5.5.

Note: If you decide to install and use your own MySQL or MariaDB server, please be aware that the class server is configured slightly differently from the default, specifically:

- default-storage-engine = INNODB (to ensure that all tables use the InnoDB storage engine)
- sql-mode = "ANSI,ONLY_FULL_GROUP_BY,STRICT_ALL_TABLES" (to ensure that students aren't nipped by MySQL's non-standard grouping/aggregation support, or its non-standard support of double-quoted strings)

MySQL Workbench
The MySQL Workbench is a very helpful graphical tool for working with your database. This is the tool that most students use to write their solutions. You can download the workbench here:
http://dev.mysql.com/downloads/tools/workbench/

Once you have downloaded this tool, you can create and save a connection to your database, so that it is very easy to connect and do your work.

---

[1] http://dev.mysql.com/doc/refman/5.5/en/
[2] http://dev.mysql.com/doc/refman/5.5/en/string-functions.html#function_lower

MySQL Console

You can also use the MySQL console client; this client is available on the CMS lab machines, and it is also available if you install MySQL or MariaDB on your own machine. The arguments to use for the console client are included in the email you will receive when your database is created. This tool isn't particularly recommended; the graphical client is much easier to use.

If you are working at the console, you can save the details of your interactions using the MySQL client's "tee" feature, and then load the results into an editor later. Just run the client as follows:

```
mysql [other arguments] --tee=hw2out.txt
```

As you use the MySQL client, every single command you type, and its output, will be stored into the file **hw2out.txt**.

Also, if you get really aggravated by the MySQL client's tendency to beep loudly at you for typing errors, look in the manual for the **no-beep** option. It's so useful that you might want to put it into **mysql.cnf**.

Setting Up Database Schemas

Once you have connected to the database server using your own username and password, you can load and run the script specified in each section to create the schema for that part of the assignment.

- If you are using the command-line client then use the **source** command. Note that this isn't an actual SQL statement; it is simply a command the client provides for importing SQL files. Use the **help** command to learn about **source** and the other client commands you can use.
- If you are using the MySQL Workbench, you can choose "Open SQL Script…" from the File menu, load the file, and then run it from the workbench.

If the file runs correctly, you should see no errors. You may see a few warnings for the very first commands; these are because the SQL schema file will try to drop tables if they already exist; if they don't exist, you get a little warning.

You can verify that these steps have completed successfully by running these simple commands:

```
SHOW TABLES;
SHOW TABLE STATUS;
```

(The MySQL Workbench has a graphical representation of the database schema, so you don't need to do this with the graphical UI.)

The **SHOW** commands are very useful for finding out what is in a database schema. Look at the **SHOW TABLE STATUS** output in particular; it shows some of the underlying details of the tables in your database. For example, we can see that the database is using the InnoDB storage engine, which provides us with both transactions and referential integrity. (The default MyISAM storage engine provides neither of these important capabilities, but it is much faster and is frequently used for loading and analyzing very large datasets in MySQL.)

Part A: Book Problems (40 points)

Do the following problems from the "Exercises" section of Chapter 3 of the textbook. These problems require you to write SQL statements to perform various queries or modifications to a database. Follow good SQL coding style.

Make sure you adhere only to MySQL-supported features. (Unfortunately, MySQL doesn't support the **WITH** clause, which is definitely helpful for some of the problems.)

The schema for the university database is in **make-university.sql**.
The schema for the library database is in **make-library.sql**.

All of your work for this problem will go into the `bookproblems.sql` file.

Problems:

1. [Exercise 3.11] *(Each part is worth 3 points; total is 9 points.)*
    Write the following queries in SQL, using the university schema (described at end of the set).

    a) Find the names of all students who have taken at least one Comp. Sci. course; make sure there are no duplicate names in the result.

    c) For each department, find the maximum salary of instructors in that department. You may assume that every department has at least one instructor.

    d) For the lowest, across all departments, of the per-department maximum salary computed by the preceding query. You should actually include the query from part c; you can't just pretend it has a name (or use a view, etc.).

2. [Exercise 3.12] *(Parts a-b are worth 2 points, c-f are worth 3 points; total is 16 points.)*

    a) Create a new course "CS-001", titled "Weekly Seminar", with 0 credits. The course ID is "CS-001", which is stored as a string.

    b) Create a section of this course in Autumn 2009, with *sec_id* of 1.

    c) Enroll every student in the Comp. Sci. department in the above section. Don't forget the **INSERT ... SELECT** statement.

    d) Delete enrollments in the above section where the student's name is Chavez. You would normally also want to update the student's *tot_cred* "total credits" value, but since CS-001 is worth 0 credits, just concentrate on removing the appropriate rows from the *takes* table, and ignore the other part.

    e) Delete the course CS-001. What will happen if you run this delete statement without first deleting offerings (sections) of this course?

    f) Delete all *takes* tuples corresponding to any section of any course with the word "database" as a part of the title; ignore case when matching the word with the title. You might find the

LOWER() function helpful. See §3.4.2 or the MySQL 5.5 manual[2] for more details on this function.

3. [Exercise 3.21] *(Part a is worth 3 points; b-d are worth 4 points; total is 15 points.)*
    Consider the following library database schema:

> *member* (*memb_no*, *name*, *age*)
> *book* (*isbn*, *title*, *authors*, *publisher*)
> *borrowed* (*memb_no*, *isbn*, *date*)

a) Retrieve the names of members who have borrowed <u>any</u> book published by "McGraw-Hill".

b) Retrieve the names of members who have borrowed <u>all</u> books published by "McGraw-Hill".

c) For each publisher, retrieve the names of members who have borrowed more than five books of that publisher.

d) Compute the average number of books borrowed per member. Take into account that if a member does not borrow any books, then that member does not appear in the *borrowed* relation at all. To be more specific, members who borrow no books should bring down the average, as you would expect, since the number of books they have borrowed is 0.

Part B: Relational Algebra and SQL (14 points)

In this problem you will translate relational algebra expressions into corresponding SQL queries. Note that these are not just fragments of SQL; you should write a <u>complete</u> SQL query that would produce the <u>exact</u> same results as the corresponding relational algebra expression.

Write your solutions in a file called `algebra.sql`.

Given two relation schemas $R = (A, B, C)$, and $S = (D, E, F)$, and two relations $r(R)$ and $s(S)$, write SQL queries that produce identical results to the following relational algebra expressions. Recall that r and s are sets of tuples, not multisets, and the results of these operations should also be sets of tuples. Therefore you should eliminate duplicates where necessary, but only do this where it's actually necessary! Feel free to explain your rationale in the comments.

1. $\Pi_A(r)$

2. $\sigma_{B=17}(r)$

3. $r \times s$

4. $\Pi_{A,F}(\sigma_{C=D}(r \times s))$

For these problems, $R = (A, B, C)$ as before, and let there be two relations $r_1(R)$ and $r_2(R)$. Write complete SQL queries (not SQL fragments) that produce identical results to these relational algebra expressions. Again, consider that the inputs are sets of tuples, and the results must also be sets of tuples.

5. $r_1 \cup r_2$

---

[2] http://dev.mysql.com/doc/refman/5.5/en/string-functions.html#function_lower

6. $r_1 \cap r_2$

7. $r_1 - r_2$

(Scoring: Each part is worth 2 points, for a total of 14 points.)

Part C: MySQL Database Exercises (46 points)

The remaining exercises in this assignment use the SQL schema for the bank database that we have discussed so far. **Note that the schema in many of the lecture slides is different from the schema in the SQL file, so review the SQL file before trying these exercises.**

All of your work for this part will be done in the `banking.sql` file.

The banking schema is provided in the file **`make-banking.sql`**. For reference, here is the banking database schema:

> *branch (branch_name, branch_city, assets)*
> *customer (customer_name, customer_street, customer_city)*
> *loan (loan_number, branch_name, amount)*
> *borrower (customer_name, loan_number)*
> *account (account_number, branch_name, balance)*
> *depositor (customer_name, account_number)*

Problem 1 *(5 parts, 2 points per part; total is 10 points)*

Here are some warm-up exercises for you. Write a SQL query for each of these questions. Make sure your query runs against the database and generates a correct result.

a) Retrieve the loan-numbers and amounts of loans with amounts of at least $1000, and at most $2000.

b) Retrieve the loan-number and amount of all loans owned by Smith. Order the results by increasing loan number.

c) Retrieve the city of the branch where account A-446 is open.

d) Retrieve the customer name, account number, branch name, and balance, of accounts owned by customers whose names start with "J". Order the results by increasing customer name.

e) Retrieve the names of all customers with more than five bank accounts.

Problem 2 *(Part a is 2 points, parts b-c are 3 points each; total is 8 points)*

Write SQL statements that define the following views. Make sure that all columns in each view's results have meaningful names and are in the order specified.

a) A view called **`pownal_customers`** containing the account numbers and customer names (but not the balances) for all accounts at the Pownal branch.

b) A view called **`onlyacct_customers`** containing the name, street, and city of all customers who have an account with the bank, but do not have a loan. This view can be defined such that it is updatable; for full points you must make sure that it is.

c) A view called **branch_deposits** that lists all branches in the bank, along with the total account balance of each branch, and the average account balance of each branch. Make sure all computed values are given reasonable names in the view.

Make sure that branches with no accounts are included in the result! (The Markham branch has no accounts.) Such branches should have a total balance of 0, and an average of **NULL**. (You may find the MySQL **IFNULL()** function useful. See the MySQL 5.1 documentation, §11.4 "Control Flow Functions," for details.)

Problem 3 *(Parts a-b are 3 points each, c-f are 4 points each; 22 points total)*

Here are some more challenging problems to try against the banking database.

a) Generate a list of all cities that customers live in, where there is no bank branch in that city. Make sure that the results are distinct; no city should appear twice. Also, sort the cities in increasing alphabetical order.

b) Are there any customers who have neither an account nor a loan? Write a SQL query that reports the name of any customers that have neither an account nor a loan. Note that MySQL does not support the **EXCEPT** operator! But there is more than one way...

c) The bank decides to promote its branches located in the city of Horseneck, so it wants to make a $50 gift-deposit into all accounts held at branches in the city of Horseneck. Write the SQL **UPDATE** command for performing this operation.

Do not use MySQL-specific extensions for this **UPDATE** operation. You may only use the standard "**UPDATE tblname SET ... WHERE ...**" form, where you can only specify one table in the **UPDATE** clause, and all references to other tables must appear in the **WHERE** clause.

d) MySQL also supports a non-standard, multiple table version of **UPDATE**, of the form:

    UPDATE tbl1, tbl2, ... SET ... WHERE ...

In this form, you can refer to columns from any of the specified tables in the **SET** and **WHERE** clauses, and MySQL will figure out what to update from what you write. Write another answer to part c, using this syntax. Note that you can also give the tables in the **UPDATE** clause aliases, as usual.

e) Retrieve all details (*account_number*, *branch_name*, *balance*) for the largest account at each branch. Implement this query as a join against a derived relation in the **FROM** clause.

f) Implement the same query as in the previous problem, this time using an **IN** predicate with multiple columns, e.g. "**(branch_name, balance) IN ...**"

Problem 4: Rank *(6 points)*

A *rank* of a collection of records assigns a 1 to the top-most value (it is in 1st place), a 2 to the second highest value (2nd place), and so forth. If multiple records have the same value, they are assigned the same rank, but there is a corresponding gap to the next rank value. For example, if you had the following values:

- 350
- 190
- 470
- 350

Their rank would be as follows:

- 1: 470
- 2: 350
- 2: 350
- 4: 190

In particular, note that no value has a rank of 3, because two values have a rank of 2.

If a value has a rank of 2, this means there is one other value that is greater. If a value has a rank of 4, there are three other values that are greater. Given this, it should be evident that you need to compare the set of records against itself in order to properly compute the rank.

Compute the rank of all bank branches, based on the amount of assets that each branch holds. The result schema should be *(branch_name, assets, rank)*, where the *rank* value is the rank of that branch.

Hint 1: There are several approaches you can use to compute this result, but all of them will come down to the same general theme: joining the branches table against itself with a non-equality condition. We call these "non-equijoins."

Hint 2: There are two branches with the same assets; make sure the rank values are correct for those branches and for the adjacent branches in the ranking.

Feedback Survey (+3 bonus points)
Complete the feedback survey for this assignment on the course website, and 3 points will be added to your score (max of 100/100).
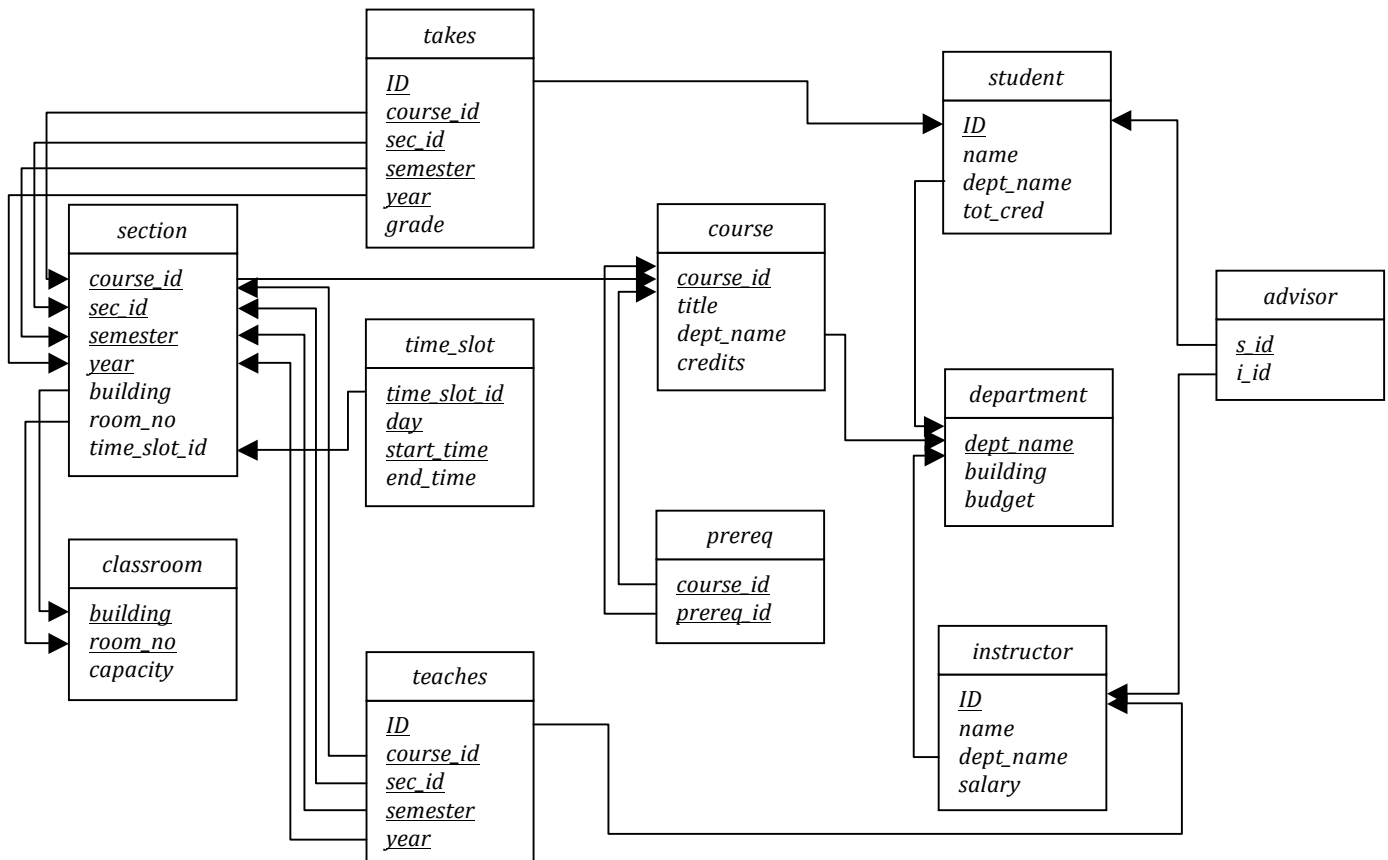
University Database

Figure 2.8 Schema diagram for the university database.

classroom (building, room_number, capacity)
department (dept_name, building, budget)
course (course_id, title, dept_name, credits)
instructor (ID, name, dept_name, salary)
section (course_id, sec_id, semester, year, building, room_number, time_slot_id)
teaches (ID, course_id, sec_id, semester, year)
student (ID, name, dept_name, tot_cred)
takes (ID, course_id, sec_id, semester, year, grade)
advisor (s_ID, i_ID)
time_slot (time_slot_id, day, start_time, end_time)
prereq (course_id, prereq_id)

Figure 2.9 Schema of the university database.

create table *department*
    (*dep_name*   **varchar** (20),
    building   **varchar** (15),
    budget   **numeric** (12,2),
    primary key (*dept_name*));

create table *course*
    (*course_id*   **varchar** (7),
    title   **varchar** (50),
    dept_name   **varchar** (20),
    credits   **numeric** (2,0),
    primary key (*course_id*),
    foreign key (*dept_name*) **references** *department*);

create table *instructor*
    (*ID*   **varchar** (5),
    name   **varchar** (20) **not null**,
    dept_name   **varchar** (20),
    salary   **numeric** (8,2),
    primary key (*ID*),
    foreign key (*dept_name*) **references** *department*);

create table *section*
    (*course_id*   **varchar** (8),
    sec_id   **varchar** (8),
    semester   **varchar** (6),
    year   **numeric** (4,0),
    building   **varchar** (15),
    room_number **varchar** (7).
    time_slot_id   **varchar** (4),
    primary key (*course_id, sec_id, semester, year*),
    foreign key (*course_id*) **references** *course*);

create table *teaches*
    (*ID*   **varchar** (5),
    course_id   **varchar** (8),
    sec_id   **varchar** (8),
    semester   **varchar** (6),
    year   **numeric** (4,0),
    primary key (*ID, course_id, sec_id, semester, year*),
    foreign key (*course_id, sec_id, semester, year*) **references** *section*,
    foreign key (*ID*) **references** *instructor*);

Figure 3.1 SQL data definition for part of the university database.